1. Basics of AI & ML

What is Artificial Intelligence?

Artificial Intelligence (AI) is a broad field of computer science focused on creating intelligent machines that can simulate human-like cognitive functions. These functions include learning, reasoning, problem-solving, perception, and natural language understanding. The ultimate goal of AI is to create systems that can perform tasks that would normally require human intelligence. Think of AI as the umbrella term for any technique that allows a machine to mimic human intelligence.

- 2. Difference between AI, Machine Learning, and Deep Learning. * AI (Artificial Intelligence) is the big-picture concept, the field of making machines intelligent. It includes everything from simple rule-based systems to complex neural networks. It's the goal.
 - ML (Machine Learning) is a subset of AI. It's a method that allows machines to learn from data without being explicitly programmed. Instead of writing rules for every possible scenario, you give the model data and it learns the patterns itself.
 - DL (Deep Learning) is a subset of ML. It uses a specific type of algorithm called an Artificial Neural Network (ANN) with multiple layers (hence "deep"). Deep learning is especially good at tasks involving large amounts of unstructured data like images, text, and audio.
- 3. Types of Machine Learning (Supervised, Unsupervised, Reinforcement).
 - Supervised Learning: The model learns from labeled data. You have input data and corresponding correct outputs. The goal is to learn a mapping from the inputs to the outputs. It's like a student learning with a teacher.
 - **Example:** A model that predicts house prices based on features like size and location, where you have a dataset of house sizes and their known prices.
 - Unsupervised Learning: The model learns from unlabeled data, meaning there are no pre-defined correct outputs. The goal is to find hidden patterns, structures, or relationships within the data. It's like a student exploring on their own.
 - Example: Grouping customers into segments based on their purchasing behavior without any prior knowledge of what those segments are.
 - Reinforcement Learning: The model (agent) learns by interacting with an environment. It performs actions and receives rewards or penalties based on the outcome. The goal is to maximize the cumulative reward over time. It's like training a pet with treats.
 - **Example:** An AI agent learning to play a game like chess or Go, where it gets a reward for winning and a penalty for losing.

4. What are common applications of AI/ML in real life?

- Image Recognition: Face unlocking on your phone, identifying objects in photos.
- Natural Language Processing (NLP): Virtual assistants like Siri and Alexa, translation services like Google Translate, and spam filters in email.
- **Recommendation Engines:** Netflix suggesting what to watch next, Amazon recommending products you might like.
- **Self-Driving Cars:** Using computer vision and other sensors to navigate roads.
- Financial Fraud Detection: Banks using ML to flag suspicious transactions.

5. What is the difference between rule-based AI and learning-based AI?

- Rule-Based AI: This is a traditional approach where developers manually write explicit rules or instructions for the system to follow. If-then-else logic is a common example. It's deterministic and predictable but can't adapt to new data or situations not covered by the rules.
 - **Example:** A chatbot that responds with pre-written answers based on keywords. If a user types "hello," it responds with "Hi, how can I help you?"
- Learning-Based AI: This is the core of modern ML. The system learns the rules and patterns directly from data. It can generalize to new, unseen data and adapt over time.
 - Example: A spam filter that learns to identify spam emails by analyzing thousands of examples, without being explicitly told what specific keywords to look for.

6. What is the difference between predictive and descriptive analytics?

- Descriptive Analytics: This answers the question, "What happened?" It summarizes and describes past data. It's about understanding trends and patterns that already exist in your dataset.
 - **Example:** Creating a report that shows the total sales for each quarter of the past year.
- **Predictive Analytics:** This answers the question, **"What will happen?"** It uses historical data and statistical models to make predictions about future outcomes.
 - **Example:** Using past sales data, customer demographics, and marketing campaigns to forecast next quarter's sales.

7. What is feature engineering? **

Feature engineering is the process of using domain knowledge to select, transform, or create new variables (features) from raw data. The goal is to create features that make the learning algorithm work better and improve model performance. It's often the most crucial step in the ML pipeline.

- **Example:** If you're building a model to predict house prices, instead of using just the raw date, you might create a new feature like "age of the house" (current year minus the build year) or "season" (spring, summer, etc.), as these new features might be more predictive.
- 8. What is dimensionality reduction?

Dimensionality reduction is the process of reducing the number of features (variables) in a dataset. This is done to combat the "curse of dimensionality," where models perform

poorly with too many features, especially if many are redundant or irrelevant.

- **Benefits:** It can reduce model complexity, prevent overfitting, and speed up training. It also helps with data visualization.
- Common Techniques: Principal Component Analysis (PCA) and t-SNE.
- Example: If you have a dataset of customer information with 1000 columns, you
 might use dimensionality reduction to combine related features and reduce the
 dataset to 50 more meaningful components.
- 9. What is bias-variance tradeoff? 1

The bias-variance tradeoff is a core concept in ML that describes the conflict between a model's ability to learn and its ability to generalize to new data.

- **Bias:** A model with **high bias** makes simplifying assumptions about the data, causing it to underfit. It's too simple and can't capture the underlying patterns.
- Variance: A model with high variance is too complex. It learns the noise and random fluctuations in the training data, leading to overfitting. It performs great on training data but poorly on new data.
- The Tradeoff: You can't have both low bias and low variance. Decreasing one usually increases the other. The goal is to find a sweet spot—a model with a good balance that generalizes well.

10. What is overfitting vs underfitting?

- Overfitting: This happens when a model learns the training data too well, including the noise and random errors. It performs exceptionally well on the training data but fails to generalize to new, unseen data. It's like memorizing the answers to a test instead of understanding the concepts.
 - How it looks: Training accuracy is high, but validation/test accuracy is low.
- Underfitting: This happens when a model is too simple to capture the underlying patterns in the training data. It performs poorly on both training and new data. It's like not studying for a test at all.
 - How it looks: Both training and validation/test accuracy are low.

2. Supervised Learning

11.

What is supervised learning? Examples.

Supervised learning is a type of ML where the model learns from a dataset that includes both the input features and the corresponding correct output labels. The goal is to learn a mapping function from the inputs to the outputs so the model can make accurate predictions on new data.

• Classification: Predicting a discrete label.

- Example: Email spam detection (Spam/Not Spam), image classification (Cat/Dog).
- **Regression:** Predicting a continuous value.
 - **Example:** Predicting house prices (\$), predicting a person's age.

12. Difference between classification and regression.

- Classification: The output variable is a category or a discrete label. The model classifies an input into one of several predefined classes.
 - Example: Predicting if a customer will churn (Yes/No), identifying the type of fruit in an image (Apple/Orange/Banana).
- **Regression:** The output variable is a **real number** or a continuous value. The model predicts a numerical value.
 - **Example:** Predicting the temperature tomorrow, predicting the price of a car.
- 13. Explain linear regression with an example. 📈

Linear regression is a simple supervised learning algorithm used for regression tasks. It assumes a linear relationship between the input variables (X) and the single output variable (Y). The goal is to find the best-fitting straight line that minimizes the sum of the squared differences between the predicted values and the actual values.

- The Equation: Y=mX+b, where m is the slope and b is the intercept. In ML terms, these are the model's coefficients.
- Example: Predicting a student's test score (Y) based on the number of hours they studied (X). The model would learn a line that shows how a student's score increases with more study hours.

14. What is logistic regression?

Logistic regression is a supervised learning algorithm used for classification tasks, despite its name. It models the probability that a given input belongs to a certain class. It's primarily used for binary classification (e.g., Yes/No, O/1). It uses a sigmoid function to squash the output of a linear equation into a probability value between 0 and 1.

Example: Predicting if an email is spam. The model outputs a probability (e.g., 0.85), which is then converted to a class (e.g., Spam) if it exceeds a certain threshold (e.g., 0.5).

15. What are decision trees? 🌳

A decision tree is a supervised learning algorithm that works by creating a tree-like model of decisions and their possible consequences. It splits the data into smaller and smaller subsets based on a series of feature-based questions, forming a hierarchical structure. Each internal node represents a test on an attribute, each branch represents an outcome of the test, and each leaf node represents a class label or a final decision.

 Example: A decision tree to decide whether to play tennis. The first question might be "Is it sunny?" If yes, "Is the humidity high?" If no, "Is the wind strong?" and so on until it reaches a final decision to Play or Not Play.

16. What is random forest?

Random Forest is an ensemble learning method that works by building a large number of decision trees and combining their predictions. The "random" part comes from two sources:

- 1. Each tree is trained on a random subset of the data (bootstrapping).
- 2. Each tree considers only a random subset of features at each split.
- How it works: For classification, the final prediction is the class with the most votes from all trees. For regression, it's the average of all tree predictions. This reduces the risk of a single tree overfitting and significantly improves overall accuracy and robustness.

17. What is gradient boosting (XGBoost, LightGBM, CatBoost)?

Gradient Boosting is another ensemble learning technique. Unlike Random Forest, which builds trees independently, gradient boosting builds trees sequentially. Each new tree is trained to correct the errors (residuals) of the previous trees. It's a powerful technique that often achieves state-of-the-art results.

- **XGBoost:** Known for its speed, performance, and advanced features like regularization. It's one of the most popular and widely-used boosting algorithms.
- LightGBM: Optimized for speed and low memory usage, especially for large datasets. It uses a different splitting strategy (histogram-based) than XGBoost.
- CatBoost: Excels at handling categorical features automatically without requiring pre-processing like one-hot encoding.

18. What are Support Vector Machines (SVM)?

Support Vector Machines (SVM) are a powerful supervised learning algorithm used for both classification and regression. For classification, the goal is to find the best hyperplane (a decision boundary) that separates the data into different classes. The "best" hyperplane is the one that has the largest margin between the closest data points of each class. These closest points are called "support vectors."

- Key Feature: SVM can use different kernels (e.g., linear, polynomial, radial basis function) to handle non-linear data by projecting it into a higher-dimensional space where it can be linearly separated.
- 19. What are k-Nearest Neighbors (k-NN)?

k-Nearest Neighbors (k-NN) is a simple, non-parametric, supervised learning algorithm that can be used for both classification and regression. It's based on the idea that similar things exist in close proximity.

- How it works: To make a prediction for a new data point, k-NN looks at the k-closest data points in the training set.
 - For Classification: The new data point is assigned the class that is most common among its k-neighbors.
 - For Regression: The new data point's value is the average of the values of its k-neighbors.
- **Note:** It's a "lazy learner" because it doesn't build a model during the training phase; it just stores the data. The computation happens during the prediction phase.

20. How to evaluate supervised learning models?

Model evaluation is crucial to understand how well a model performs on new data.

- Classification Metrics:
 - **Accuracy:** The ratio of correctly predicted instances to the total instances. Good for balanced datasets.

- **Precision:** Of all the positive predictions, how many were actually correct?
- Recall (Sensitivity): Of all the actual positive instances, how many did the model correctly identify?
- **F1-Score:** The harmonic mean of precision and recall. Useful when you need a balance between them, especially with imbalanced classes.
- ROC-AUC (Receiver Operating Characteristic Area Under Curve): A plot of the true positive rate vs. the false positive rate. AUC measures the model's ability to distinguish between classes. A value closer to 1 is better.

Regression Metrics:

- Mean Absolute Error (MAE): The average of the absolute differences between the predicted and actual values.
- Mean Squared Error (MSE): The average of the squared differences. Penalizes large errors more.
- **R-squared (R2):** Represents the proportion of the variance in the dependent variable that can be explained by the independent variables.

3. Unsupervised Learning

21. _

What is unsupervised learning? Examples. 🕵

Unsupervised learning is a type of ML where the model is given unlabeled data and must find hidden patterns or structures within it. There's no "correct" output to guide the algorithm.

- **Clustering:** Grouping similar data points together.
 - **Example:** Segmenting customers into groups based on their purchasing habits to create targeted marketing campaigns.
- **Dimensionality Reduction:** Reducing the number of features in a dataset.
 - **Example:** Compressing image data while retaining key information.
- Association Rule Mining: Finding relationships between items.
 - **Example:** The "people who bought this also bought..." feature on Amazon.
- 22. What is clustering? Explain K-Means.

Clustering is the task of grouping a set of objects in such a way that objects in the same group (cluster) are more similar to each other than to those in other groups.

- **K-Means:** A popular, simple clustering algorithm.
 - 1. Choose a number of clusters, k.
 - 2. Randomly initialize k cluster centroids.
 - 3. Assign each data point to the nearest centroid.
 - 4. Update the centroid's position to be the mean of all points assigned to that

cluster.

5. Repeat steps 3 and 4 until the centroids no longer move significantly.

23. Difference between K-Means and Hierarchical clustering.

- K-Means: Requires you to pre-specify the number of clusters (k). It's a "partitioning" method. It's computationally efficient and good for large datasets. The clusters are often spherical.
- **Hierarchical Clustering:** Does not require you to specify the number of clusters beforehand. It builds a hierarchy of clusters. There are two types:
 - Agglomerative (bottom-up): Starts with each data point as its own cluster and merges the closest pairs until all points are in one cluster.
 - **Divisive (top-down):** Starts with all data points in one cluster and recursively splits them.
 - **Pros:** Can reveal a nested structure in the data. **Cons:** More computationally expensive than K-Means and less suitable for very large datasets.

24. What is DBSCAN clustering?

DBSCAN (Density-Based Spatial Clustering of Applications with Noise) is a clustering algorithm that groups together points that are closely packed together, marking as outliers (noise) points that lie alone in low-density regions.

- **Key Idea:** It finds clusters based on density rather than distance from a centroid (like K-Means). It can find clusters of arbitrary shapes and is robust to outliers.
- Core Concepts:
 - Core Point: A point with at least a minimum number of neighbors within a certain radius.
 - **Border Point:** A point with fewer neighbors than the minimum but is in the neighborhood of a core point.
 - Noise Point: A point that is not a core point and is not a border point.

25. What is PCA (Principal Component Analysis)?

PCA is an unsupervised dimensionality reduction technique. Its goal is to transform a high-dimensional dataset into a lower-dimensional one while retaining as much of the original variance as possible.

How it works: It identifies new, uncorrelated axes (called principal components)
that capture the most variance in the data. The first principal component captures
the most variance, the second captures the second most, and so on. You can then
project the data onto a smaller number of these components.

26. Difference between PCA and t-SNE.

- PCA (Principal Component Analysis):
 - Goal: Dimensionality reduction while preserving the global structure (variance) of the data.
 - Output: New, uncorrelated features (principal components).
 - Use Case: Often used as a pre-processing step for ML algorithms to reduce features.
- t-SNE (t-Distributed Stochastic Neighbor Embedding):
 - Goal: Dimensionality reduction primarily for visualization. It preserves the local

structure of the data.

- Output: A 2D or 3D map where similar data points are clustered together.
- Use Case: Excellent for visualizing high-dimensional data, like showing clusters of similar documents or images, but not suitable as a pre-processing step for ML models.

27. What is anomaly detection?

Anomaly detection (or outlier detection) is the task of identifying rare items, events, or observations that deviate significantly from the majority of the data. These anomalies can be an indication of a problem or a new opportunity.

- Applications:
 - Credit card fraud detection: A transaction that is unusually large or from a strange location.
 - **Network security:** Detecting unusual network traffic that might indicate a cyber attack.
 - Manufacturing: Identifying defective products on an assembly line.
- 28. What is recommendation system (collaborative vs content-based)?

A recommendation system is a type of filtering system that predicts user preferences and recommends items they might like.

- Collaborative Filtering: Recommends items based on the behavior of similar users.
 It looks for patterns in user-item interactions.
 - **Example:** "Users who liked this movie also liked these movies." It doesn't need information about the items themselves.
- **Content-Based Filtering:** Recommends items that are similar to items the user has liked in the past. It uses features of the items.
 - **Example:** If a user likes action movies, the system will recommend other action movies based on genre, actors, and other content features.
- 29. What is dimensionality reduction and why is it used? (Duplicate question, but rephrased for context)

Dimensionality reduction is the process of reducing the number of variables (dimensions) in a dataset. It is used for several key reasons:

- Curse of Dimensionality: In high-dimensional spaces, data becomes sparse, making it harder for models to find patterns.
- Overfitting: A large number of features can lead to models that are too complex and overfit the training data.
- **Improved Performance:** Reducing the number of features can speed up model training and improve generalization.
- **Data Visualization:** It's easier to visualize data when the number of dimensions is reduced to 2 or 3.
- 30. How do you evaluate clustering models?

Evaluating unsupervised models is tricky because there are no labels. You evaluate based on how "good" the clusters are.

 Silhouette Score: Measures how similar a data point is to its own cluster compared to other clusters. The score ranges from -1 to 1. A high score means the data point is

- well-matched to its own cluster and poorly matched to neighboring clusters.
- Davies-Bouldin Index: Measures the average similarity between clusters. A lower score indicates better clustering. It's calculated by taking the ratio of the intra-cluster distance to the inter-cluster distance. A value of 0 indicates a perfect separation.

4. Reinforcement Learning

31.

What is reinforcement learning?

Reinforcement Learning (RL) is a type of ML where an agent learns to make decisions by taking actions in an environment to maximize a cumulative reward. The agent is not told what to do; instead, it must discover the best actions through trial and error.

 Analogy: Teaching a dog a new trick. The dog (agent) performs actions (sits, stands, barks). You (the environment) provide a reward (a treat) when it performs the desired action. The dog learns to associate the action with the reward.

32. Explain agent, environment, state, action, reward.

- Agent: The learner or decision-maker. It takes actions in the environment.
- **Environment:** The world the agent interacts with. It provides the current state and responds to the agent's actions.
- **State:** The current situation or snapshot of the environment.
- **Action:** A move the agent makes in the environment.
- **Reward:** A feedback signal (positive or negative) from the environment in response to an action. The agent's goal is to maximize the cumulative reward.

33. Difference between model-based and model-free RL.

- Model-Based RL: The agent learns a model of the environment. This model predicts the next state and reward for each possible action. The agent can then use this model to plan ahead and choose the best actions.
 - **Example:** A chess AI that learns the rules of the game and then simulates different moves to find the best one.
- Model-Free RL: The agent does not learn a model of the environment. Instead, it learns directly from trial and error by taking actions and observing the rewards. It learns a policy (a strategy) for which action to take in each state.
 - **Example:** Q-Learning and Policy Gradients. The agent learns the value of each action in each state without understanding the underlying dynamics of the environment.

34. What is Q-Learning?

Q-Learning is a popular model-free reinforcement learning algorithm that learns an action-value function, denoted as Q(s,a). The "Q" stands for "quality."

- Goal: To learn the value of taking a certain action (a) in a particular state (s). A higher Q-value means a better action.
- How it works: It updates the Q-values based on the rewards received after each action, using an update rule called the Bellman equation. The agent then chooses the action with the highest Q-value in a given state.

35. What is Deep Q-Learning (DQN)?

Deep Q-Learning (DQN) is an extension of Q-Learning that uses a deep neural network to approximate the Q-function.

- Problem with traditional Q-Learning: It uses a large table to store Q-values, which
 is not feasible for complex environments with huge numbers of states (e.g., video
 games).
- DQN's Solution: A neural network takes the state as input and outputs the Q-value for each possible action. This allows it to generalize from a limited number of experiences and handle high-dimensional state spaces like images.

36. What is Policy Gradient?

Policy Gradient is a type of RL algorithm that directly learns a policy function (π) that maps states to actions. Instead of learning the value of each action in each state (like Q-Learning), it learns a direct strategy.

 How it works: It uses gradient-based optimization to adjust the parameters of the policy function to increase the probability of taking actions that lead to higher rewards.

37. Difference between on-policy and off-policy learning.

- On-Policy Learning: The agent learns the optimal policy by evaluating and improving the same policy that it uses to act (to make decisions). The policy is being constantly updated.
 - **Example:** SARSA (State-Action-Reward-State-Action). It evaluates the value of the action based on the next action that the current policy will take.
- Off-Policy Learning: The agent learns an optimal policy by evaluating a policy that
 is different from the one it uses to act. It can learn from data collected by other
 policies or from past experiences.
 - Example: Q-Learning. It learns the value of the optimal action (the one with the highest Q-value), regardless of whether the current policy would actually choose that action.

38. What is exploration vs exploitation?

This is a fundamental tradeoff in RL.

- **Exploitation:** The agent chooses the action that it currently believes will yield the highest reward. It's like a chef using a recipe they know works well.
- **Exploration:** The agent tries a new, unproven action to discover if it might lead to an even higher reward. It's like a chef trying a new ingredient to see if it improves the dish.
- The Tradeoff: An agent needs to balance these two. Too much exploitation leads to a sub-optimal policy (getting stuck in a local maximum). Too much exploration leads to slow learning and might miss out on rewards.

39. Examples of reinforcement learning in real life.

- o Gaming: Google's AlphaGo and AlphaZero, which mastered Go and chess.
- **Robotics:** A robot learning to navigate a room or perform a task like grasping an object.
- **Finance:** Algorithmic trading systems that learn to buy or sell stocks to maximize profit.
- Resource Management: Optimizing energy consumption in data centers.
- 40. What are Markov Decision Processes (MDP)?

A Markov Decision Process (MDP) is a mathematical framework for modeling sequential decision-making in environments where outcomes are partly random and partly under the control of a decision-maker. It's the theoretical foundation of reinforcement learning.

 Key Property: The Markov Property, which states that the future is independent of the past given the present state. In simpler terms, the current state contains all the information needed to make an optimal decision; you don't need to know the history of how you got there.

5. Deep Learning

41. _

What is deep learning? How is it different from ML?

Deep Learning (DL) is a subset of ML that uses Artificial Neural Networks (ANNs) with multiple layers to learn from data.

- How it's different from traditional ML:
 - Feature Engineering: In traditional ML, you often manually create features.

 Deep learning models can automatically learn and extract features from raw data
 - **Data Scale:** DL requires and performs best with very large datasets. Traditional ML can work with smaller datasets.
 - **Performance:** DL models often outperform traditional ML models on complex tasks like image and speech recognition.
 - Black Box: DL models are often considered "black boxes" because their decision-making process is hard to interpret. Traditional ML models are generally more transparent.
- 42. What are artificial neural networks (ANN)? 🧠

An Artificial Neural Network (ANN) is a computing system inspired by the structure of the human brain. It consists of interconnected nodes (neurons) organized in layers:

- Input Layer: Takes the raw data.
- Hidden Layers: One or more layers that perform computations and find patterns.

This is where the "deep" in deep learning comes from.

- o **Output Layer:** Produces the final result.
- How it works: Neurons in one layer are connected to neurons in the next. Each connection has a weight. As data passes through the network, each neuron applies an activation function to the weighted sum of its inputs, and the output is passed to the next layer. The network "learns" by adjusting these weights.

43. Explain perceptron and multi-layer perceptron.

- Perceptron: The simplest form of a neural network. It's a single-layer feed-forward network used for binary classification. It takes multiple inputs, calculates a weighted sum, and applies a step function to produce a binary output (0 or 1). It can only solve linearly separable problems.
- Multi-Layer Perceptron (MLP): A more advanced and widely-used type of neural network. It consists of an input layer, one or more hidden layers, and an output layer. The hidden layers allow the MLP to solve non-linearly separable problems, which a simple perceptron cannot. This is because the hidden layers can learn complex, non-linear relationships.
- 44. What is activation function? Examples (ReLU, Sigmoid, Tanh, Softmax).

An activation function is a crucial component of a neural network neuron. It introduces non-linearity into the model, allowing it to learn complex patterns and relationships. Without activation functions, a neural network would just be a linear model, no matter how many layers it has.

- ReLU (Rectified Linear Unit): The most common activation function. It outputs the input directly if it's positive, and 0 otherwise. Pros: Simple, computationally efficient.
- **Sigmoid:** Squashes the input to a range between 0 and 1. **Pros:** Good for binary classification outputs. **Cons:** Can suffer from the vanishing gradient problem.
- Tanh (Hyperbolic Tangent): Squashes the input to a range between -1 and 1. Pros: Zero-centered, which can help with optimization.
- **Softmax:** Used in the output layer for **multi-class classification**. It converts a vector of numbers into a probability distribution, where the sum of all outputs is 1.

45. What is forward propagation and backpropagation?

- o **Forward Propagation:** The process of feeding input data through the neural network to get an output. Data flows from the input layer, through the hidden layers, to the output layer. During this process, the network makes a prediction.
- Backpropagation: The core algorithm for training a neural network. After a forward pass, the model's prediction is compared to the actual value to calculate the error (loss). This error is then propagated backward from the output layer to the input layer. During this backward pass, the algorithm calculates the gradient of the loss with respect to each weight, and then updates the weights to reduce the error in the next forward pass.
- 46. What are convolutional neural networks (CNN)?

Convolutional Neural Networks (CNNs) are a class of deep learning models specifically designed for processing data with a grid-like topology, such as images.

• Key Feature: They use a mathematical operation called convolution to learn

features from the data. A **kernel** (filter) slides over the input data, creating a **feature map**. This process allows the network to automatically learn hierarchical features, from simple lines and edges to more complex shapes and objects.

Typical Layers:

- Convolutional Layer: Applies filters to the input.
- Pooling Layer: Down-samples the feature map to reduce its size and complexity.
- Fully Connected Layer: Performs classification based on the extracted features.

Image of a convolutional neural network architecture

47. Applications of CNN (image recognition, object detection).

- Image Recognition (Classification): Classifying an entire image into a single category (e.g., Is this a picture of a cat or a dog?).
- Object Detection: Identifying and locating multiple objects within an image and drawing a bounding box around them (e.g., Finding all the cars, people, and traffic lights in a photo).
- Other Applications: Facial recognition, medical image analysis (detecting tumors),
 video analysis, and satellite imagery analysis.

48. What are recurrent neural networks (RNN)? ...

Recurrent Neural Networks (RNNs) are a class of neural networks designed for processing sequential data, such as time series, text, and audio.

- Key Feature: They have a "memory" or an internal state that allows them to process sequences of arbitrary length by maintaining information from previous steps. They do this by feeding the output of a layer back into the input of the same layer for the next time step.
- Limitations: Simple RNNs can struggle with the vanishing gradient problem, making it difficult to learn long-term dependencies.

49. Difference between RNN, LSTM, and GRU.

- RNN (Recurrent Neural Network): Basic RNNs struggle to remember information over long sequences due to the vanishing gradient problem. The influence of early inputs fades as the sequence grows.
- LSTM (Long Short-Term Memory): A special type of RNN designed to overcome
 the vanishing gradient problem. LSTMs have a more complex internal structure with
 "gates" (input, forget, output) that control the flow of information. This allows them
 to selectively remember or forget information, enabling them to learn long-term
 dependencies.
- GRU (Gated Recurrent Unit): A simpler variant of LSTM. It has only two gates (reset and update), making it computationally more efficient than LSTMs while still performing well on many tasks. GRUs are a good alternative when computational resources are a concern.

50. What is transformer architecture?

The Transformer architecture is a revolutionary neural network architecture, introduced in the "Attention Is All You Need" paper, that has become the dominant model for NLP

tasks.

- Key Innovation: It completely discards the recurrent (RNN) and convolutional (CNN) layers and relies entirely on a mechanism called self-attention.
- How it works: The self-attention mechanism allows the model to weigh the importance of different words in a sequence when processing a single word. This allows it to capture long-range dependencies in the text much more efficiently than RNNs.
- **Examples:** BERT, GPT, and T5 are all based on the Transformer architecture.

6. NLP (Natural Language Processing)

51. ____

What is NLP?

Natural Language Processing (NLP) is a field of AI focused on enabling computers to understand, interpret, and generate human language. It involves various tasks, from simple text processing to complex language generation.

52. Explain Bag-of-Words model.

The Bag-of-Words (BoW) model is a simple way to represent text data for machine learning. It converts text into a numerical vector by counting the frequency of each word in a document.

- How it works:
 - 1. Create a vocabulary of all unique words in the corpus.
 - 2. For each document, create a vector of the same length as the vocabulary.
 - 3. The value at each position in the vector is the count of how many times that word appears in the document.
- **Limitation:** It completely ignores word order and context, treating each word as an independent feature.

53. What is TF-IDF?

TF-IDF (Term Frequency-Inverse Document Frequency) is an improvement over the Bag-of-Words model. It's a numerical statistic that reflects how important a word is to a document in a collection or corpus.

- **Term Frequency (TF):** How often a word appears in a document.
- Inverse Document Frequency (IDF): A measure of how rare or unique a word is across the entire corpus. Words that appear in many documents (like "the" or "a") have a low IDF.
- **TF-IDF Calculation:** TF-IDF=TF×IDF. This formula gives more weight to words that are frequent in a specific document but rare in the overall corpus, making them more significant.

54. What is word2vec?

Word2vec is a technique for generating word embeddings. It learns to represent words as dense, low-dimensional vectors in a way that captures their semantic and syntactic relationships.

 Key Idea: It's based on the distributional hypothesis: "a word is known by the company it keeps." Words that appear in similar contexts will have similar vector representations.

Two Architectures:

- **Skip-gram:** Predicts the surrounding context words given a target word.
- CBOW (Continuous Bag of Words): Predicts a target word from the surrounding context words.
- Cool Feature: It can perform vector arithmetic. For example, vector("king")-vector("man")+vector("woman")≈vector("queen").

55. Difference between word2vec, GloVe, and fastText.

- Word2vec: Learns embeddings by predicting words from their context (Skip-gram) or context from the word (CBOW). It's a predictive model.
- GloVe (Global Vectors for Word Representation): Combines the ideas of word2vec (local context) with global matrix factorization. It uses a co-occurrence matrix (how often words appear together) to learn the embeddings.
- **fastText:** An extension of word2vec. It treats each word as a combination of its sub-word units (character n-grams). This allows it to:
 - Handle out-of-vocabulary (OOV) words by creating an embedding from its character n-grams.
 - Generate embeddings for rare words more effectively.

56. What are embeddings in NLP?

Embeddings are a technique used to represent categorical data, like words, as a numerical vector of real numbers. These vectors are typically dense (all elements are non-zero) and low-dimensional.

Why they're important: They capture semantic relationships between words. Words
with similar meanings or contexts will be located close to each other in the vector
space. This allows ML models to understand the relationships between words, which
is a significant improvement over simple one-hot encoding.

57. What is Named Entity Recognition (NER)?

Named Entity Recognition (NER) is an NLP task that identifies and classifies named entities in text into predefined categories, such as person names, organizations, locations, dates, and so on.

 Example: In the sentence "Steve Jobs worked at Apple in Cupertino," an NER model would identify "Steve Jobs" as a Person, "Apple" as an Organization, and "Cupertino" as a Location.

58. What is sentiment analysis?

Sentiment analysis (or opinion mining) is the NLP task of determining the sentiment or emotional tone expressed in a piece of text.

• **Goal:** To classify the text as positive, negative, or neutral.

- **Applications:** Analyzing customer reviews, social media feedback, or market research data to understand public opinion.
- 59. What are attention mechanisms?

An attention mechanism is a technique used in deep learning models, particularly in NLP, that allows the model to weigh the importance of different parts of the input sequence.

- How it works: When a model is processing a word, the attention mechanism calculates a score for how relevant every other word in the input sequence is to the current word. It then uses these scores to create a weighted sum of all the words, which is used to make a better prediction. This helps the model focus on the most relevant parts of the input, especially for long sequences.
- 60. What are large language models (LLMs)? Examples (GPT, BERT, T5).

Large Language Models (LLMs) are a type of transformer-based deep learning model that has been trained on a massive amount of text and code data.

- Key Characteristics:
 - Scale: They have a huge number of parameters (billions or even trillions).
 - **Pre-training:** They are first trained on a vast and diverse corpus of data in a self-supervised manner (e.g., predicting the next word).
 - Generative vs. Discriminative:
 - **Generative (e.g., GPT):** Can generate new, coherent text.
 - **Discriminative (e.g., BERT):** Can understand and classify existing text.
- Examples:
 - GPT (Generative Pre-trained Transformer): Developed by OpenAI, known for its ability to generate human-like text.
 - BERT (Bidirectional Encoder Representations from Transformers): Developed by Google, known for its ability to understand the context of words in both directions.
 - T5 (Text-to-Text Transfer Transformer): Developed by Google, which frames all NLP tasks as a text-to-text problem.

7. Model Training & Evaluation

and a second

61. -

What is cross-validation? 🧪

Cross-validation is a technique used to evaluate a model's performance and prevent overfitting. It involves splitting the dataset into multiple subsets or "folds."

 Why it's used: Instead of a single train-test split, which can be sensitive to the specific data points in the test set, cross-validation provides a more robust and reliable estimate of a model's performance on unseen data.

62. What is k-fold cross-validation?

k-fold cross-validation is the most common form of cross-validation.

Steps:

- 1. The dataset is divided into \mathbf{k} equal-sized folds.
- 2. The model is trained k times.
- 3. In each iteration, one fold is used as the validation set, and the remaining k-1 folds are used as the training set.
- 4. The performance metric is calculated for each iteration.
- 5. The final performance is the average of the k metrics. A common value for k is 5 or 10.

63. What is stratified sampling?

Stratified sampling is a technique used when splitting a dataset into training, validation, and testing sets, particularly for classification problems with imbalanced classes.

- **Goal:** To ensure that the proportion of each class in the splits (train/test) is the same as in the original dataset.
- **Example:** If your dataset for a fraud detection model has 99% non-fraudulent transactions and 1% fraudulent ones, a standard random split might result in a test set with no fraudulent transactions. Stratified sampling would ensure that the test set also has 1% fraudulent transactions, giving you a more realistic evaluation.

64. What is confusion matrix?

A confusion matrix is a table used to evaluate the performance of a classification model. It summarizes the number of correct and incorrect predictions made by the model.

Components:

- True Positive (TP): Predicted positive, actual positive.
- False Positive (FP): Predicted positive, actual negative. (Type I Error)
- True Negative (TN): Predicted negative, actual negative.
- False Negative (FN): Predicted negative, actual positive. (Type II Error)

65. Explain precision, recall, and F1-score.

These metrics are derived from the confusion matrix and are crucial for evaluating classification models, especially with imbalanced datasets.

- Precision: Precision=TP+FPTP. Of all the positive predictions the model made, how many were correct? It measures the model's accuracy when predicting the positive class.
- Recall (Sensitivity): Recall=TP+FNTP. Of all the actual positive instances, how many did the model correctly identify? It measures the model's ability to find all the positive cases.
- F1-Score: F1-Score=2×Precision+RecallPrecision×Recall. The harmonic mean of precision and recall. It's a good metric to use when you need a balance between precision and recall, as it penalizes models with extremely poor performance in one of the two metrics.

66. What is ROC curve and AUC?

 ROC Curve (Receiver Operating Characteristic Curve): A graph that shows the performance of a classification model at all classification thresholds. It plots the True Positive Rate (Recall) against the False Positive Rate (FPR).

- AUC (Area Under the Curve): The area under the ROC curve. It provides a single number that summarizes the model's performance. An AUC of 1.0 is a perfect classifier, and 0.5 is a random classifier. A higher AUC means the model is better at distinguishing between the positive and negative classes.
- 67. What is regularization? Difference between L1 and L2.

Regularization is a technique used to prevent overfitting in machine learning models by adding a penalty to the loss function. This penalty discourages the model from becoming too complex and learning the noise in the training data.

- L1 Regularization (Lasso): Adds a penalty equal to the absolute value of the magnitude of the coefficients. It can force some coefficients to become exactly zero, effectively performing feature selection.
- L2 Regularization (Ridge): Adds a penalty equal to the squared magnitude of the coefficients. It forces the coefficients to be small but rarely makes them exactly zero. It helps with multicollinearity and makes the model more robust.
- 68. What is dropout in neural networks?

Dropout is a regularization technique specific to neural networks. During training, it randomly "drops out" (sets to zero) a certain percentage of neurons in a layer.

- Why it works: It prevents neurons from co-adapting too much. By randomly dropping out neurons, the network is forced to find multiple independent representations, making the model more robust and less likely to overfit. It's like training a collection of different neural networks at once.
- 69. What is gradient descent? Variants (SGD, Adam, RMSprop).

Gradient Descent is an optimization algorithm used to find the minimum of a function, typically the loss function in an ML model. It works by iteratively moving in the direction opposite to the gradient of the function.

- SGD (Stochastic Gradient Descent): Updates the weights for each training example, one at a time. It's computationally faster but has a noisy path to the minimum.
- Adam (Adaptive Moment Estimation): A more advanced optimizer that combines
 the best features of other optimizers. It's often the default choice. It adapts the
 learning rate for each parameter and is very efficient.
- o **RMSprop (Root Mean Square Propagation):** Similar to Adam, it adapts the learning rate for each parameter but doesn't use momentum like Adam. It is effective for non-stationary objectives (e.g., in RNNs).
- 70. What is hyperparameter tuning? (Grid search, Random search, Bayesian optimization). Hyperparameter tuning is the process of finding the optimal set of hyperparameters for a model to achieve the best performance. Hyperparameters are the settings for the algorithm itself (e.g., learning rate, number of hidden layers, number of trees in a random forest) and are not learned from the data.
 - **Grid Search:** Exhaustively searches a pre-defined subset of hyperparameters. It's simple but computationally expensive.
 - o Random Search: Selects random combinations of hyperparameters. It's more

- efficient than Grid Search, as it often finds a good combination in fewer iterations.
- Bayesian Optimization: A more intelligent approach. It uses past results to inform the selection of the next set of hyperparameters, finding the optimal values much more efficiently.

8. Model Deployment & Scaling

71. –

How to deploy ML models in production? 🚀

Deploying an ML model means making its predictions available to other applications or users.

- Common Steps:
 - 1. **Serialization:** Save the trained model in a standard format (e.g., pickle, ONNX).
 - 2. **API Creation:** Wrap the model in a web service (e.g., using Flask or FastAPI) that exposes an API endpoint to make predictions.
 - 3. **Containerization:** Package the application and all its dependencies into a container (e.g., Docker).
 - 4. **Deployment:** Deploy the container to a cloud platform (e.g., AWS, Azure, Google Cloud) or an on-premise server.
 - 5. **Monitoring:** Set up monitoring to track the model's performance and health.

72. Difference between batch inference and real-time inference.

- Batch Inference: Making predictions on a large batch of data all at once, usually on a scheduled basis (e.g., daily or hourly). The data is static, and the predictions are stored for later use.
 - **Example:** Running a model nightly to predict which customers are likely to churn.
- Real-Time Inference (Online Inference): Making predictions on a single data point or a small number of data points as they arrive. It requires low latency.
 - **Example:** A recommendation system on an e-commerce website that makes a product recommendation as a user browses.

73. What is model monitoring (drift detection)?

Model monitoring is the process of tracking the performance and health of a deployed model over time.

- Drift Detection: A key part of monitoring. Models trained on historical data can become less accurate over time because the underlying data distribution changes. This is called model drift or data drift.
- Example: A model trained to predict product demand based on seasonal trends might become less accurate if consumer behavior shifts unexpectedly. Monitoring helps detect this drift so the model can be retrained.

74. What is A/B testing in ML models?

A/B testing (or split testing) is a method used to compare two versions of an ML model (Model A and Model B) to determine which one performs better in a production environment.

 How it works: A portion of the incoming traffic or users is routed to Model A, while another portion is routed to Model B. You then measure a key metric (e.g., click-through rate, conversion rate) to see which model provides a better outcome.

75. How to scale ML models (sharding, distributed training)?

- **Sharding:** A database technique applied to ML where you break up a large dataset into smaller, more manageable pieces (shards) and process them in parallel.
- **Distributed Training:** Training a model across multiple machines or GPUs to handle very large datasets or complex models that a single machine cannot handle.
 - Data Parallelism: Each machine gets a different chunk of the data but uses the same model.
 - Model Parallelism: A single, large model is split across multiple machines, with each machine holding a part of the model.

76. What is transfer learning?

Transfer learning is a technique where a model trained on one task is re-purposed or adapted for a different but related task.

- Why it's used: It's very effective when you have a small dataset for your target task, as you can leverage the knowledge gained from a large, pre-trained model. It saves a significant amount of time and computational resources.
- **Example:** Using a pre-trained image recognition model (e.g., VGG, ResNet) that was trained on millions of images and fine-tuning it with a small dataset to classify a specific type of image, like different species of flowers.

77. What is fine-tuning in ML models?

Fine-tuning is a specific method of transfer learning. After a model has been pre-trained on a large dataset, you continue to train it on a smaller, task-specific dataset.

 Process: The model's final layers are often removed and replaced with new ones relevant to the new task. The entire network (or just the new layers) is then trained for a few epochs on the new data, with a very small learning rate, so the pre-trained weights are only slightly adjusted.

78. What is model explainability (SHAP, LIME)? 🧐

Model explainability is the field of techniques that aims to make the predictions of ML models understandable to humans. This is especially important for complex "black box" models like deep neural networks.

- SHAP (SHapley Additive exPlanations): A game theory-based approach that explains a prediction by calculating the contribution of each feature to the prediction.
- LIME (Local Interpretable Model-agnostic Explanations): Explains the predictions
 of any black box model by approximating it with a simple, interpretable model (like a
 linear model or decision tree) in the local vicinity of the prediction.

79. What is MLOps?

MLOps (Machine Learning Operations) is a set of practices that combines Machine Learning, DevOps, and Data Engineering to create a seamless, end-to-end process for deploying and maintaining ML models in production.

 Goal: To bridge the gap between model development (by data scientists) and model deployment (by engineers) by automating the entire ML lifecycle, including data collection, model training, testing, deployment, and monitoring.

80. What are ML pipelines?

An ML pipeline is a sequence of steps that automates the entire machine learning workflow.

- Typical Steps: Data ingestion, data cleaning, feature engineering, model training, model evaluation, and model deployment.
- Benefits:
 - **Automation:** Reduces manual effort and potential for human error.
 - **Reproducibility:** Ensures that the same results can be achieved every time.
 - Scalability: Allows you to handle larger datasets and more complex models.
 - Collaboration: Provides a standardized workflow for teams.

9. Advanced Topics

81.

What is generative AI? Examples (GANs, Diffusion models, LLMs). Generative AI is a type of AI that can create new, original data (images, text, audio) that is similar to the data it was trained on but is not a copy of it.

- **How it works:** These models learn the underlying patterns and structure of the training data and use this knowledge to generate new content.
- Examples:
 - GANs (Generative Adversarial Networks): A generator model creates new data, and a discriminator model tries to determine if it's real or fake. They train in a competitive game.
 - **Diffusion Models:** A more recent class of models that learn to create images by progressively removing noise from a pure noise image.
 - LLMs (Large Language Models): Can generate new text, code, and other forms of content.
- 82. What are GANs (Generative Adversarial Networks)?

A GAN consists of two neural networks, a Generator and a Discriminator, that are trained simultaneously in a zero-sum game.

• **The Generator:** Its job is to generate new data (e.g., images) that are as realistic as possible to fool the discriminator.

- The Discriminator: Its job is to distinguish between real data from the training set and fake data generated by the generator.
- The Training Process: The generator gets better at creating realistic data, while the
 discriminator gets better at spotting the fakes. They push each other to improve until
 the generator can create data that is indistinguishable from real data.

Image of a generative adversarial network

83. Difference between GAN and Variational Autoencoder (VAE).

o GANs:

- **Goal:** To generate high-quality, realistic data.
- **Training:** Adversarial training (two networks competing).
- Output: Often produces very sharp and realistic images.

VAEs (Variational Autoencoders):

- Goal: To learn a compressed representation (latent space) of the data and generate new samples from it.
- **Training:** Uses a single neural network with an encoder and a decoder. It's a probabilistic approach.
- Output: Tends to produce blurrier or less sharp images than GANs but offers a more structured latent space, which is useful for tasks like interpolation.

84. What are diffusion models? (like Stable Diffusion).

Diffusion models are a new class of generative AI models that have recently become very popular for their ability to generate high-quality images.

• **Core Idea:** They work by adding noise to an image and then learning to reverse that process to create a new image from scratch (starting with pure noise).

Process:

- 1. **Forward Diffusion:** An image is progressively corrupted by adding Gaussian noise over several steps.
- 2. **Reverse Diffusion:** A neural network (often a U-Net) is trained to predict and subtract the noise at each step, essentially learning to "denoise" the image until it becomes a clear, new image.
- **Advantage:** They can generate extremely high-quality and diverse images. Stable Diffusion is a well-known example.

85. What is reinforcement learning with human feedback (RLHF)?

Reinforcement Learning from Human Feedback (RLHF) is a method used to align large language models with human values and preferences.

How it works:

- 1. An LLM generates multiple responses to a prompt.
- 2. Human annotators rank or rate these responses based on quality, helpfulness, or safety.
- 3. A **reward model** is trained on this human feedback to predict which responses humans would prefer.
- 4. The LLM is then fine-tuned using reinforcement learning, with the reward model as the guide, to produce responses that maximize the predicted human

preference.

86. What are foundation models?

A foundation model is a large-scale, pre-trained model that can be adapted for a wide range of downstream tasks.

- Key Idea: It learns foundational knowledge and representations from a vast and diverse dataset, and this knowledge can be transferred to new tasks with minimal fine-tuning.
- Examples: GPT-4 and BERT. They are trained on a massive amount of text from the internet, and then can be fine-tuned for tasks like sentiment analysis, text summarization, or question answering.

87. What is multi-modal AI?

Multi-modal AI refers to AI systems that can process and understand data from multiple modalities (types of data) at the same time.

- **Common Modalities:** Text, images, audio, and video.
- Examples:
 - A model that can generate an image from a text description (e.g., DALL-E, Midjourney).
 - A model that can answer questions about an image (e.g., "What is the dog doing in this picture?").

88. What are vector databases?

A vector database is a database designed to store, manage, and search vector embeddings.

- Why they are needed: Traditional databases are not efficient for searching through high-dimensional vector data. Vector databases use specialized indexes (e.g., approximate nearest neighbor search) to find similar vectors very quickly.
- Application: They are essential for applications like semantic search (finding documents with similar meaning), recommendation systems, and RAG (Retrieval-Augmented Generation) in LLMs.

89. Difference between symbolic AI and neural networks.

- Symbolic AI (Good Old-Fashioned AI): This is a traditional approach that uses symbols, logic, and explicit rules to represent knowledge and solve problems. It's based on human-defined rules.
 - Example: Expert systems, where a system is given a set of if-then rules to diagnose a problem.
 - **Pros:** Interpretable, explainable. **Cons:** Brittle, not good at handling ambiguity or large amounts of unstructured data.
- **Neural Networks:** This is a modern approach that learns patterns and representations directly from data. It's a "black box" approach.
 - **Pros:** Excellent at handling complex, unstructured data; can adapt and generalize. **Cons:** Less interpretable.

90. What is federated learning?

Federated learning is a decentralized machine learning approach that trains an algorithm across multiple distributed devices (e.g., mobile phones, hospitals) without the need to

centralize the data.

- How it works:
 - 1. A central server sends a model to the devices.
 - 2. Each device trains the model on its local data.
 - 3. The devices send back the **model updates** (weights and gradients), not the raw data, to the central server.
 - 4. The server aggregates these updates to create a new, improved model.
- Benefits: It enhances privacy and security because the data never leaves the device.

10. Real-World ML

91.

What are common challenges in ML projects?

- Data Quality: Messy, incomplete, or biased data is a major challenge. "Garbage in, garbage out."
- o Data Scarcity: Not having enough labeled data to train a good model.
- **Feature Engineering:** This is a manual and time-consuming process.
- **Model Overfitting:** Creating a model that performs well on training data but poorly in production.
- Model Explainability: Understanding why a complex model made a certain prediction.
- Model Deployment & MLOps: Getting a model from development to production and maintaining it.
- 92. What is data leakage in ML? 💧

Data leakage is a serious problem that occurs when information from the test or validation set "leaks" into the training process. This causes a model to perform unrealistically well on the training data, but it will fail in a real-world setting.

- Example: A model for credit card fraud detection accidentally uses a feature like transaction_was_marked_as_fraudulent in the training data, which would not be available at the time of prediction.
- 93. What is class imbalance? How to handle it?

Class imbalance occurs in classification problems when the number of data points for one class is significantly lower than for another.

- **Example:** In a fraud detection dataset, fraudulent transactions might make up only 1% of the data.
- How to Handle It:
 - Resampling:

- Oversampling: Duplicating examples from the minority class (e.g., using SMOTE).
- Undersampling: Removing examples from the majority class.
- Using different metrics: Instead of accuracy, use precision, recall, or F1-score, which are more informative for imbalanced data.
- Cost-sensitive learning: Assign a higher penalty to misclassifying the minority class.

94. What is cold start problem in recommendation systems?

The cold start problem occurs when a recommendation system cannot make accurate recommendations because it lacks sufficient data. This happens for two main reasons:

- **New User:** The system doesn't know the new user's preferences because they haven't interacted with enough items yet.
- **New Item:** A new item is added to the system, but no one has interacted with it yet, so the system can't recommend it.
- **Solutions:** Use content-based filtering for new users (e.g., ask them for their interests) or popular items for new items.

95. What is active learning in ML?

Active learning is a machine learning approach where a model can interactively query an expert (e.g., a human annotator) to label new data points.

- Goal: To achieve high performance with a minimum amount of labeled training data.
 The model intelligently selects the most informative, uncertain, or representative data points to ask for a label.
- **Use Case:** When getting labeled data is expensive or time-consuming.

96. What is semi-supervised learning?

Semi-supervised learning is a hybrid approach that uses both a small amount of labeled data and a large amount of unlabeled data for training.

• **How it works:** The model can use the unlabeled data to learn the underlying structure of the data and use the labeled data to guide the learning process. It's particularly useful when getting a lot of labeled data is difficult.

97. What is self-supervised learning?

Self-supervised learning is a method where a model learns from a vast amount of unlabeled data by creating its own supervisory signals.

- **How it works:** A model is trained on a "pretext task" where the labels are automatically generated from the data itself.
- Example: An image model might be trained to predict the color of a grayscale image or to predict a hidden part of an image. The model learns a rich representation of the data that can then be used for a separate, downstream task.

98. What is continual learning?

Continual learning (or lifelong learning) is a type of machine learning where the model can learn from a continuous stream of new data without forgetting what it has learned from previous data.

o Challenge: Most models suffer from catastrophic forgetting, where training on new data overwrites the knowledge gained from old data. Continual learning aims to

solve this problem.

• **Applications:** Robots learning to perform new tasks, or a model that needs to be updated with new information as it becomes available.

99. What is ethical AI?

Ethical AI is a field that focuses on the responsible development and use of artificial intelligence. It addresses the potential for AI systems to cause harm, whether intentionally or unintentionally.

- Key Concerns:
 - Bias: All models can reflect and amplify biases present in their training data (e.g., a hiring All that shows bias against certain genders).
 - Fairness: Ensuring that AI systems do not discriminate against protected groups.
 - Accountability: Who is responsible when an AI makes a harmful decision?
 - **Transparency:** The ability to understand how an AI system makes its decisions.

100. Future of AI/ML – where is it going? 🚀

The field of AI is evolving rapidly, with several key trends:

- * Scaling Up: Models are getting larger, with more parameters and are trained on more data, leading to a phenomenon known as "emergent abilities."
- * Multi-Modality: Al systems that can seamlessly integrate and understand text, images, video, and audio are becoming the norm.
- * Generative AI: The rise of models like GANs, diffusion models, and LLMs is revolutionizing content creation and many other industries.
- * Edge AI: Moving AI processing from the cloud to the device itself (e.g., on your phone or in a car) to improve privacy, latency, and reliability.
- * Responsible AI: There's a growing focus on building AI systems that are fair, transparent, and trustworthy.
- * Personalization: AI models will become even more tailored to individual users, from personalized medicine to hyper-specific content recommendations.